

Abdullah

To implement the three variants of QuickSort I used a generalized implementation of the in-place QuickSort that takes in via lambda expression a function that picks the pivot. I went with this approach as it meant less code duplication, easier debugging, and easier maintenance. It handles all edge cases, and will throw an exception when given a null array. For partitioning I used a modified version of the Hoare partition scheme which places the pivot as the last element. I chose Hoare partitioning as it had linear time and the only other option I knew of was Lomuto partitioning and that had exponential time.

Pivoting on the first element was a trivial implementation and it completes the operation in one step not affecting the time complexity as it is an $O(1)$ operation.

The same held true for pivoting on a random element as I used the built in `java.util.Random` implementation to help in picking the pivot, which is $O(1)$

For the Median of three implementations of QuickSort I had two parallel arrays which each contain three elements, one contains a value and the other the corresponding index. The arrays are then sorted with a simple bubble sort according to the values array. Now one would think this would ruin the run time, but it does not as this bubble sort only ever sorts three elements per pass, making it in reality a $O(1)$ operation as it is not affected by the size of the input directly

The implementations match the theoretical performance of $O(n\log(n))$ for most cases but when it came to the worst case of a reverse ordered array the pivot on first experienced $O(n^2)$ even failing due to a stack overflow for the largest array test run. Whilst median of three did the best in the reverse ordered array, pivoting on random took slightly longer than the median of three.

Abenezer

Bucket Sort's implementation employs dynamic bucket sizing based on the square root of the array length, balancing memory usage with performance. The code handles edge cases effectively, including empty arrays, single-element arrays, and identical element arrays. Using `ArrayList` for buckets provides efficient memory management for variable bucket sizes. The implementation normalizes value ranges to properly handle negative numbers and prevent division by zero errors. While alternative implementations using fixed bucket counts or `LinkedList` structures were considered, they were rejected due to poor performance with non-uniform distributions and inferior cache behavior respectively. The implementation achieves the theoretical $O(n + k)$ average case complexity for uniform distributions, while maintaining $O(n^2)$ worst case when elements cluster in a single bucket. Performance varies significantly based on data distribution - excelling with uniform data but degrading with clustered inputs.

Selection Sort uses a straightforward in-place implementation focused on minimal space usage and code simplicity. The algorithm performs a single swap per outer loop iteration to

minimize write operations. More complex variations like double-ended selection sort or recursive implementations were rejected to maintain simplicity and avoid additional space overhead. The implementation exactly matches its theoretical $O(n^2)$ time and $O(1)$ space complexity bounds. It performs a consistent $n(n-1)/2$ comparison regardless of input characteristics, making it highly predictable but inefficient for large datasets. The algorithm shows no sensitivity to input ordering - performing identically on sorted, reverse sorted, or random data.

Shell Sort implements a simple gap sequence (dividing by 2 each time) to achieve predictable though not optimal performance. The in-place implementation maintains space efficiency while moving elements within gaps similar to insertion sort. More sophisticated gap sequences like sedgwick's were considered but rejected to maintain implementation simplicity. The code achieves the theoretical $O(n^2)$ worst case bound for the chosen gap sequence, though a better sequence could theoretically approach $O(n \log n)$ at the cost of a simple sequence chosen providing predictable but sub-optimal performance. The algorithm shows particular efficiency with partially sorted arrays while handling reverse sorted arrays competently due to its gap approach.

Armando

The implementation for the Radix Sort algorithm was inspired by a YouTube channel called Geekific. I went with this implementation because the channel does a good job of breaking down complex algorithms into understandable bits with visualizations. This implementation also takes into account negative values. As for expectations, Radix Sort is considered to be better than other comparison-based sorting algorithms for large datasets when keys have several digits, although it is not the best for small datasets. Its Big O is $d * (n + b)$ where d is the number of digits, n is the number of elements, and b is the base of the number system being used. Based on my charts for this algorithm, the slope for the curve as the input grew became less steep for all scenarios. This confirms that it does in fact do better for larger datasets, whereas the smaller input sizes yielded steeper curves. I noticed that the inputs that were fifty percent ordered and random had runtime spikes for the smaller input sizes, which further validates this algorithm's inefficiency for smaller unordered data sets.

The implementation for the Heap Sort algorithm was also inspired by a YouTube channel. I went with their implementation because they visualized the way in which a heap structure is actually used to sort some dataset. Heap sort is known for having a time complexity of $n * \log(n)$ for all cases, meaning it can handle large datasets relatively well. The n comes from the height of the heap. My charts for this algorithm did not completely align with the $n * \log(n)$ runtime expectation with all array types and it instead demonstrated more of a $\log(n)$ runtime in some cases. The reversed and random input arrays seemed to influence the $\log(n)$ runtime whereas the others aligned more with the $n * \log(n)$ runtime where the slope of the curve gets gradually steeper as the size increases.

The implementation for the Bubble Sort algorithm was inspired by an instructor's code from a CS class. I went with their implementation because they had studied the algorithm and found their implementation to be a valid way of doing it. Bubble sort is considered to be inefficient when the elements are arranged in decreasing order, with a worst-case runtime of n^2 . The charts for this algorithm confirmed this expectation since the slope for the curve increased drastically as the input got larger, which aligns with an n^2 runtime. The algorithm performed somewhat sporadically for the ordered array when the size was below 4,000, but it performed as expected for all other array types.

Joe

Counting sort ended up being a little bit faster overall than merge sort did. This matches the theoretical results since counting sort is $O(n + k)$ and merge sort is $O(n \log n)$. I'm going to compare those two run times for an array size of 32,768. I will use 40,000 as the range for counting sort below since that is the highest number that will be used in our experiment.

Counting sort, $k = 40,000$, $n + k = 32,768 + 40,000 = 72,768$ operations

Merge sort, $\log n = \log 32768 = 15$, so $n \log n = 32,768 * 15 = 491,520$

From this, for array sizes not larger than 32,768 and element ranges not larger than 40,000 counting sort is the clear winner. The problem with counting sort is that it is $O(n + k)$. So, to figure out where counting sort will start being less efficient we can look at the numbers above. If we subtract n from the merge sort operations then we get $(491,520 - 32,768 = 458,752)$.

So at an integer range of 458,753 merge sort will start running faster. If we had an extremely large number range of let's say 1 billion then counting sort would be much slower. I will show those calculations below.

$n = 32,768$

$\log n = 15$

$k = 1,000,000,000$

Merge sort $= 32,768 * 15 = 491,520$

Counting sort $= 32,768 + 1,000,000,000 = 1,000,032,768$

Thus, in that experiment merge sort is running just shy of a half million operations and counting sort is running a little over 1 billion. This matches the theoretical analysis since the integer range makes no difference to merge sort but is critical in the analysis of counting sort.

Insertion sort ran as expected since it runs at $O(n^2)$. For all of the experiments merge sort's highest run time was 4ms, counting sort's highest run time was 3ms and insertion sort's highest run time was 800 ms. In small data sets of 2,048 and below insertion sort was somewhat comparable

to merge sort and insertion sort, but after that it really started to slow down. At an array size of 32,768 insertion sort will run in 32768^2 .

$$32768^2 = 1,073,741,824$$

From these results we can see that if our number range was 1 billion that at an array size of 32,768 insertions sort and counting sort almost have the same number of operations. So if we were to have an integer range of 2 billion then insertion sort would have roughly half the operations of counting sort.

With that number range, merge sort would be unchanged with its 481,520 operations.

Thus up to a number range of 458,752 counting sort is the clear winner, with a number range above that merge sort takes over and with a range of 2 billion insertion sort is better than counting sort. Thus, if the range is below 458,753 use counting sort and anything higher use merge sort. Insertion sort will always run more slowly than both except with extremely high integer ranges, but that is irrelevant because you just use merge sort in that case anyway.

For my implementations. I implemented insertion sort from memory on how the algorithm operates and didn't choose any other implementation since there is no way to get a runtime faster than quadratic time. For merge sort I chose the algorithm from the "Coding with John" youtube channel as I found this the easiest to follow and found other implementations to have unnecessary complications. For counting sort, I followed along with the explanation from the "geeksforgeeks" website and wrote my own code.

For integer arrays with a range below $\frac{1}{2}$ million I would choose counting sort and from the 3 algorithms I was assigned, after a range above $\frac{1}{2}$ million I would choose merge sort. For an embedded system with potentially less memory you would need to look at the space complexity of the given algorithms. Counting sort has a space complexity of $O(k)$ and merge sort is $O(n)$. If you are working with integers and the integer range is smaller than n , then counting sort would be the clear winner if memory is a concern.

As far as runtime variability, counting sort and merge sort weren't really affected by the different array states. Insertion sort was much slower with a reversed array, fastest with an ordered array and somewhat consistent with the others.